

Guide to IRT Invariance Tests in R

Dr. Adam W. Meade, NC State University

2020-09-09

Getting Started

To get started, there are several packages required to use this guide. If these packages are not already installed, or if you haven't updated in a while, first run the following code:

```
install.packages("psych")
install.packages("lessR")
install.packages("mirt")
```

If these packages are already installed and updated, skip the above and just load these libraries.

```
rm(list=ls()) # clear out old junk
library("psych")
library("lessR")
library('mirt')
```

Optional Function

This step is completely optional, but I have written a function to help make it easier to identify differentially functioning (DF) and non-DF items. Note, older versions of MIRT used List format to provide results of DIF tests whereas newer version use a dataframe. Let's go ahead and load it for later:

```
get.dif.items <- function(f.data,p.val=.05,parms){
  r.warnings = ""
  keep.vars <- c("X2", "df", "p") # just keep these variables
  f.data <- f.data[keep.vars]
  f.data$p = round(f.data$p,3)
  if(missing(f.data)) return('Missing model output out.list')
  f.data$sig <- ifelse(f.data$p < p.val,'dif','no_dif')
  if(!missing(parms)){
    if(nrow(f.data) == nrow(parms)){
      f.data <- cbind(f.data,parms)
    }else{
      r.warnings = "There number of item parameters doesn't match the number of items "
      r.warnings = paste(r.warnings,"given to get.dif.items. Item parameters omitted.")
    }
  }
  dif.items <- subset(f.data, sig == 'dif')
  no.dif.items <- subset(f.data, sig == 'no_dif')
  if(!missing(parms) && nrow(f.data) == nrow(parms)){
    if(nrow(no.dif.items)>1){
```

```

    no.dif.items <- no.dif.items[order(-no.dif.items$a1),]
  }
}

r.list <- list(dif_items = dif.items, no_dif = no.dif.items, warnings = r.warnings)
return(r.list)
}

```

Create data for illustration

Next we will create some data just to illustrate the process. When analyzing your own data, you can skip this step.

```

make.data <- function(N){
  set.seed(12345)
  a <- matrix(abs(rnorm(15,1,.3)), ncol=1)
  d <- matrix(rnorm(15,0,.7),ncol=1)
  d1 <- d2 <- cbind(d, d-1, d-2) # b parameters for both groups
  d2[13:15, ] <- d1[13:15, ] + 1 # here is the DIF
  itemtype <- rep('graded', nrow(a))
  dataset1 <- simdata(a, d1, N, itemtype)
  dataset2 <- simdata(a, d2, N, itemtype)
  dat <- rbind(dataset1, dataset2)
  return(dat)
}

N <- 1000
dat <- make.data(N)
group <- c(rep('Ref', N), rep('Foc', N))
focal.data <- dat[1:1000,]
ref.data <- dat[1001:2000,]
rm(N)

```

Here we have specified our sample size per group and given the groups labels. In general, you want your reference group to be the first group.

Step 1: Testing Assumptions

Parallel Analysis

Start by examining dimensionality. I'm using parallel analysis from the Psych package. I like to do this for each group separately.

```
##### check dimensionality #####  
fa.parallel(focal.data)  
fa.parallel(ref.data)
```

Im hiding the output here, but it's clearly unidimensional.

Model Fit

Next, compute stats and examine plots. I'm surpressing output here again, but you should certainly inspect your own data.

```
##### check model fit #####  
mirtCluster(4) # speed up processing  
foc.model <- mirt(focal.data, model = 1, itemtype = "graded", SE=TRUE)  
M2(foc.model)  
ref.model <- mirt(ref.data, model = 1, itemtype = "graded", SE=TRUE)  
M2(ref.model)  
(foc.fit <- itemfit(foc.model))  
(ref.fit <- itemfit(ref.model))  
  
### optional plotting  
plots.foc <- list()  
plots.ref <- list()  
for(i in 1:ncol(dat)){  
  plots.foc[[i]]<-itemfit(foc.model,empirical.plot = i)  
  plots.ref[[i]]<-itemfit(ref.model,empirical.plot = i)  
}  
plots.foc  
plots.ref
```

Check for Sufficient Number of Category Responses

With Likert-type scales, often there will be very few respondents endorsing some of the extreme response options. If you haven't already done so, consider collapsing response categories where there are fewer than 20 respondents and/or where the SE associated with the item b parameter is greater than around .4. Different items can have different numbers of response options, but the number of response options must be the same across groups. Note - these are unpublished rules of thumb, so use at your own risk.

```
apply(ref.data, 2, table)
```

```
##   Item_1 Item_2 Item_3 Item_4 Item_5 Item_6 Item_7 Item_8 Item_9 Item_10
## 0    398    611    573    315    481    367    320    593    694    752
## 1    189    170    174    218    199    238    183    207    148    123
## 2    178    108    124    207    146    194    175    105     81     69
## 3    235    111    129    260    174    201    322     95     77     56
##   Item_11 Item_12 Item_13 Item_14 Item_15
## 0     272     545     236     241     313
## 1     188     178     186     190     215
## 2     207     117     191     201     200
## 3     333     160     387     368     272
```

```
apply(focal.data, 2, table)
```

```
##   Item_1 Item_2 Item_3 Item_4 Item_5 Item_6 Item_7 Item_8 Item_9 Item_10
## 0    395    632    539    327    474    346    317    569    713    726
## 1    186    146    209    206    170    242    186    208    144    154
## 2    170    121    136    207    152    204    184    117     90     71
## 3    249    101    116    260    204    208    313    106     53     49
##   Item_11 Item_12 Item_13 Item_14 Item_15
## 0     261     544     416     408     513
## 1     174     173     209     176     223
## 2     199     115     157     192     145
## 3     366     168     218     224     119
```

Step 2: Get Item Parameters (optional)

Freely Estimated Model

This step is also optional, but usually you'll want to look at your item parameters and put them in your manuscript somewhere. Once you start doing invariance tests, you will be constraining parameters to be the same for some items. This is a good time to go ahead and get unconstrained item parameters for your manuscript.

```
model.free <- multipleGroup(dat, 1, group)
```

```
coef(model.free, simplify = TRUE)      # for the manuscript
```

```
## $Foc
## $items
##           a1      d1      d2      d3
## Item_1  1.237  0.533 -0.454 -1.505
## Item_2  1.246 -0.585 -1.624 -2.600
## Item_3  1.114 -0.370 -1.346 -2.322
## Item_4  0.897  0.902 -0.158 -1.214
## Item_5  1.120  0.098 -0.939 -1.914
## Item_6  0.378  0.565 -0.440 -1.422
## Item_7  1.193  0.966 -0.019 -0.960
## Item_8  0.958 -0.458 -1.645 -2.606
## Item_9  0.932 -0.967 -1.943 -2.828
## Item_10 0.665 -1.212 -2.100 -3.010
## Item_11 1.102  1.214  0.184 -0.879
## Item_12 1.342 -0.240 -1.271 -2.163
## Item_13 1.350  1.570  0.437 -0.616
## Item_14 1.116  1.408  0.329 -0.686
## Item_15 0.683  0.876 -0.113 -1.078
##
## $means
## F1
## 0
##
## $cov
##      F1
## F1  1
##
##
## $Ref
## $items
##           a1      d1      d2      d3
## Item_1  1.076  0.523 -0.401 -1.343
## Item_2  1.168 -0.690 -1.581 -2.679
## Item_3  0.936 -0.192 -1.280 -2.336
## Item_4  0.865  0.826 -0.164 -1.208
## Item_5  1.175  0.139 -0.751 -1.707
## Item_6  0.540  0.681 -0.374 -1.412
## Item_7  1.289  1.006 -0.025 -1.040
## Item_8  0.870 -0.324 -1.428 -2.396
## Item_9  0.873 -1.042 -2.028 -3.204
## Item_10 0.766 -1.087 -2.199 -3.228
```

```
## Item_11 0.871 1.198 0.304 -0.635
## Item_12 1.529 -0.261 -1.319 -2.217
## Item_13 1.257 0.438 -0.682 -1.661
## Item_14 1.066 0.469 -0.400 -1.496
## Item_15 0.837 -0.060 -1.169 -2.247
##
## $means
## F1
## 0
##
## $cov
##      F1
## F1 1
```

Step 3: Baseline Model

We will use likelihood ratio tests (LRTs). To start, we need to create a baseline model in which all items have parameters constrained across items. We also will take a look at the parameters in the constrained model. I'll suppress the output here, but we are also getting item parameters constrained to be equal across groups.

```
##### Baseline Model #####
model.constrained <- multipleGroup(dat, 1, group,
  invariance = c(colnames(dat), 'free_means', 'free_var'))
(constrained.parameters <- coef(model.constrained,simplify = TRUE)[[1]][[1]])
```

Step 4: First Round of LRTs

First, test each of the items by freeing the parameters of each item, one at a time. The other items serve as the anchor items (i.e., the all-others-as-anchors model).

```
##### First round of DIF analysesb - All Others As Anchors #####
(dif.drop <- DIF(model.constrained, c('a1','d1','d2','d3'), scheme = 'drop', seq_stat = .05))
```

##		AIC	AICc	SABIC	HQ	BIC	X2	df	p
##	Item_1	-4.618	-4.075	5.078	3.609	17.786	12.618	4	0.013
##	Item_2	3.201	3.743	12.896	11.427	25.604	4.799	4	0.309
##	Item_3	-7.819	-7.277	1.876	0.407	14.585	15.819	4	0.003
##	Item_4	4.543	5.085	14.238	12.769	26.947	3.457	4	0.484
##	Item_5	-10.774	-10.232	-1.079	-2.548	11.629	18.774	4	0.001
##	Item_6	1.191	1.734	10.887	9.417	23.595	6.809	4	0.146
##	Item_7	0.436	0.978	10.131	8.662	22.839	7.564	4	0.109
##	Item_8	-4.639	-4.097	5.056	3.587	17.764	12.639	4	0.013
##	Item_9	3.398	3.941	13.094	11.624	25.802	4.602	4	0.331
##	Item_10	-1.962	-1.420	7.733	6.264	20.441	9.962	4	0.041
##	Item_11	-14.171	-13.629	-4.476	-5.945	8.233	22.171	4	0.000
##	Item_12	-3.361	-2.819	6.334	4.865	19.042	11.361	4	0.023
##	Item_13	-102.548	-102.006	-92.852	-94.322	-80.144	110.548	4	0.000
##	Item_14	-57.339	-56.797	-47.644	-49.113	-34.935	65.339	4	0.000
##	Item_15	-104.997	-104.455	-95.302	-96.771	-82.593	112.997	4	0.000

```
## use the optional function to table the output
```

```
get.dif.items(f.data=dif.drop,p.val=.05,parms=constrained.parameters)
```

```
## $dif_items
##
```

##		X2	df	p	sig	a1	d1	d2	d3
##	Item_1	12.617505	4	0.013	dif	1.1232715	0.6298066	-0.320057304	-1.3095206
##	Item_3	15.818965	4	0.003	dif	0.9922609	-0.1809792	-1.209024123	-2.2204892
##	Item_5	18.774232	4	0.001	dif	1.1031775	0.2197826	-0.734651243	-1.6890229
##	Item_8	12.639474	4	0.013	dif	0.8762771	-0.3046022	-1.440640155	-2.4009223
##	Item_10	9.962481	4	0.041	dif	0.6934429	-1.0798484	-2.075703890	-3.0409324
##	Item_11	22.170960	4	0.000	dif	0.9444870	1.2822498	0.332114956	-0.6575168
##	Item_12	11.361130	4	0.023	dif	1.3980824	-0.1167109	-1.155452094	-2.0455828
##	Item_13	110.547872	4	0.000	dif	1.2719896	1.0762762	-0.005587984	-0.9825108
##	Item_14	65.338954	4	0.000	dif	1.0961640	1.0153050	0.066249884	-0.9696149
##	Item_15	112.997023	4	0.000	dif	0.7576384	0.4724196	-0.532885307	-1.5044003

```
##
## $no_dif
##
```

##		X2	df	p	sig	a1	d1	d2	d3
----	--	----	----	---	-----	----	----	----	----

```
## Item_7 7.564119 4 0.109 no_dif 1.2134407 1.0972210 0.09394862 -0.8810099
## Item_2 4.799253 4 0.309 no_dif 1.1903195 -0.5260877 -1.49005357 -2.5256397
## Item_9 4.601646 4 0.331 no_dif 0.8922566 -0.9205046 -1.90027658 -2.9154160
## Item_4 3.456962 4 0.484 no_dif 0.8671231 0.9450889 -0.07923670 -1.1286627
## Item_6 6.808646 4 0.146 no_dif 0.4449100 0.6626897 -0.36459751 -1.3718975
##
## $warnings
## [1] ""
```

Step 5: Specify a New Baseline Model using Anchor Items

We will use the A5 method from Meade and Wright (2012) in which we will choose five anchor items with the largest A parameters. Note that the `get.dif.items` function will sort non-dif items by the A parameter if supplied. I have manually specified the items that are anchors based on the output above. Specifically, I am specifying Items 2, 4, 6, 7, and 9 based on their A parameters and non-DIF initial results.

```
itemnames <- colnames(dat)
anc.items.names <- itemnames[c(2,4,6,7,9)]
test.items <- c(1,3,5,8,10:15)
model_anchor <- multipleGroup(dat, model = 1, group = group,
  invariance = c(anc.items.names, 'free_means', 'free_var'))
(anchor.parms <- coef(model_anchor, simplify = TRUE)[[1]][[1]])
```

The output here is not very interesting so we will proceed with the final round of invariance tests.

Step 6: Run the Final Invariance Tests

```
(dif.anchor <- DIF(model_anchor, c('a1','d1','d2','d3'), items2test = test.items, plotdif = TRUE))
```

```
##### Final round of DIF analyses #####
```

```
## use the optional function to table the output
```

```
get.dif.items(f.data=dif.anchor,p.val=.05)
```

```
## $dif_items
```

```
##           X2 df p sig
```

```
## Item_13 103.851 4 0 dif
```

```
## Item_14  71.323 4 0 dif
```

```
## Item_15 119.647 4 0 dif
```

```
##
```

```
## $no_dif
```

```
##           X2 df      p    sig
```

```
## Item_1   2.862 4 0.581 no_dif
```

```
## Item_3   6.604 4 0.158 no_dif
```

```
## Item_5   6.268 4 0.180 no_dif
```

```
## Item_8   4.042 4 0.400 no_dif
```

```
## Item_10  6.352 4 0.174 no_dif
```

```
## Item_11  8.568 4 0.073 no_dif
```

```
## Item_12  1.816 4 0.770 no_dif
```

```
##
```

```
## $warnings
```

```
## [1] ""
```

Using our custom function makes it easy to see that the correct items were identified as DIF items (Items 13-15)

Step 7: Compute Effect Sizes

The last step is to compute effect size estimates, as described in Meade (2010).

Test-Level Effect Sizes

```
empirical_ES(model_anchor, DIF=FALSE) # test-level effect sizes
```

```
##          Effect Size      Value
## 1          STDS -1.1113795
## 2          UTDS  2.0559964
## 3          UETSDS  1.1113795
## 4          ETSSD -0.1630801
## 5      Starks.DTFR -1.0828614
## 6          UDTFR  2.0227532
## 7          UETSDN  1.0828614
## 8 theta.of.max.test.D  0.2529304
## 9          Test.Dmax -1.2979990
```

Item-Level Effect Sizes


```
empirical_ES(model_anchor) # item-level effect sizes
```

```
##          SIDS  UIDS  SIDN  UIDN  ESSD theta.of.max.D  max.D mean.ES.foc
## item.1  0.052  0.062  0.049  0.062  0.090      -1.031  0.103      1.266
## item.2  0.000  0.000  0.000  0.000  0.000       0.478  0.000      0.674
## item.3  0.046  0.070  0.041  0.071  0.105       2.588 -0.178      0.815
## item.4  0.000  0.000  0.000  0.000  0.000       0.478  0.000      1.389
## item.5  0.097  0.097  0.097  0.097  0.180       1.370  0.173      1.073
## item.6  0.000  0.000  0.000  0.000  0.000       0.478  0.000      1.246
## item.7  0.000  0.000  0.000  0.000  0.000       0.478  0.000      1.476
## item.8  0.084  0.084  0.082  0.082  0.227       0.127  0.096      0.750
## item.9  0.000  0.000  0.000  0.000  0.000       0.478  0.000      0.493
## item.10 0.022  0.027  0.023  0.029  0.100       2.588  0.155      0.434
## item.11 0.102  0.114  0.099  0.115  0.200      -1.644  0.232      1.672
## item.12 0.032  0.057  0.035  0.059  0.053       1.685  0.177      0.879
## item.13 -0.552  0.552 -0.535  0.535 -0.900       0.202 -0.660      1.158
## item.14 -0.436  0.436 -0.425  0.425 -0.798       0.122 -0.495      1.231
## item.15 -0.557  0.557 -0.548  0.548 -1.487       0.013 -0.597      0.861
##          mean.ES.ref
## item.1          1.214
## item.2          0.674
## item.3          0.769
## item.4          1.389
## item.5          0.976
## item.6          1.246
## item.7          1.476
## item.8          0.666
## item.9          0.493
## item.10         0.412
## item.11         1.570
```

```
## item.12      0.847
## item.13      1.710
## item.14      1.667
## item.15      1.418
```

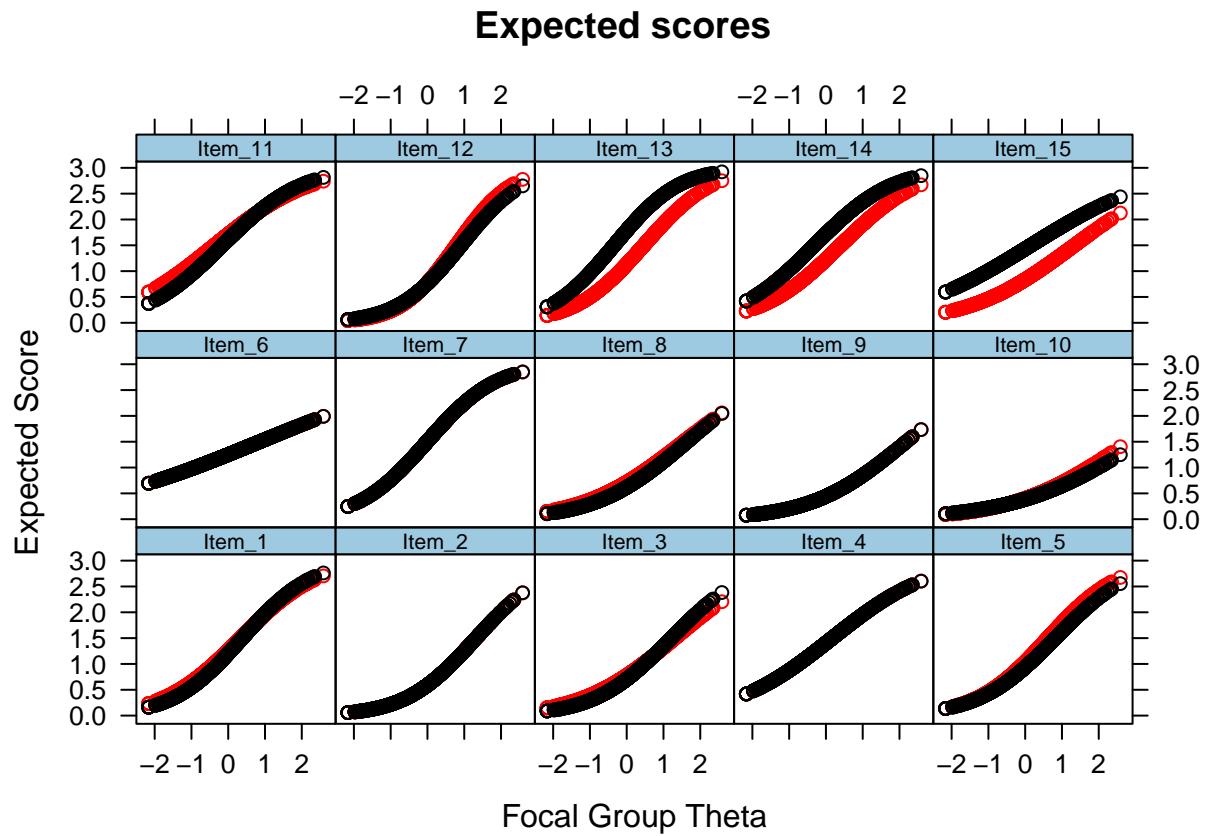
Test-Level Plot

```
expected.test.plot <- empirical_ES(model_anchor, DIF=FALSE, plot=TRUE)
expected.test.plot
```



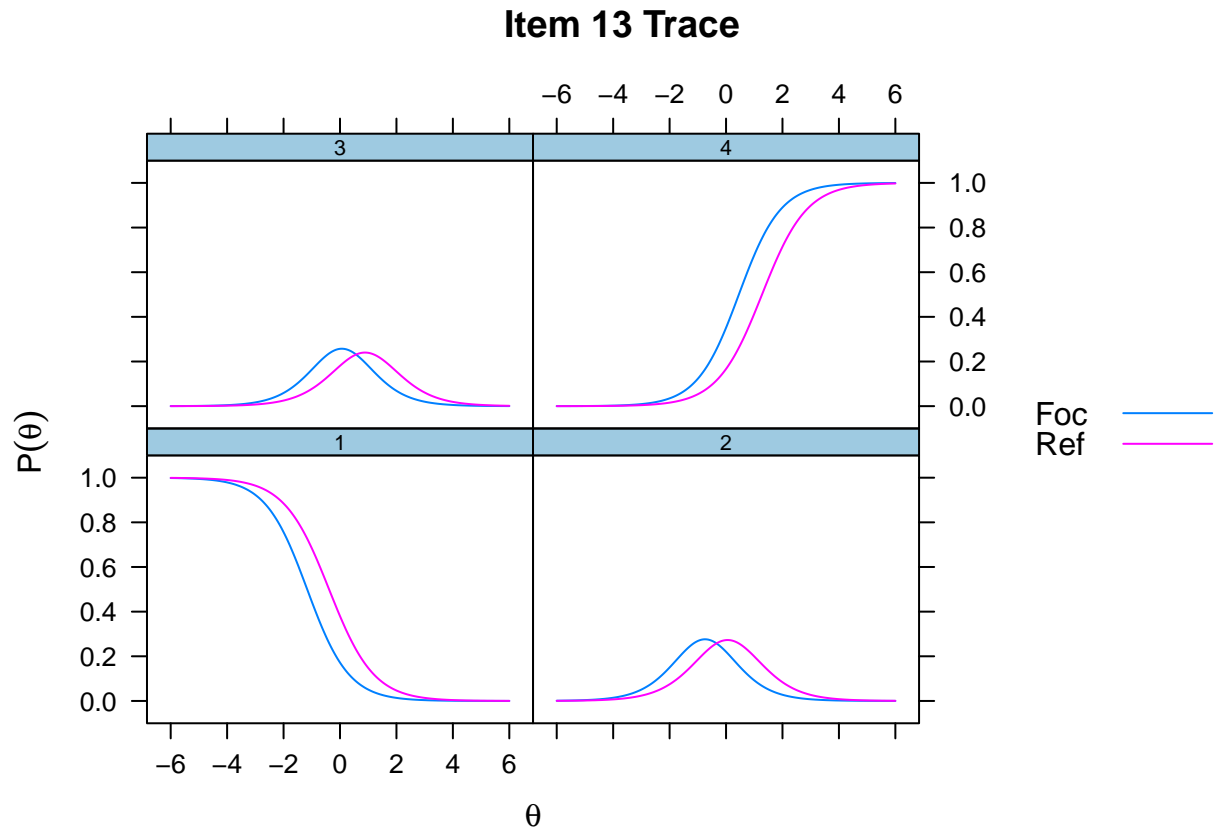
Item-Level Plots You should also examine item-level DIF plots

```
expected.item.plots <- empirical_ES(model_anchor, plot=TRUE)  
expected.item.plots
```



If you want to plot only a single item in more detail, you can do so like this

```
itemplot(model_anchor, 13) # further investigate item with DF
```



Plots are lattice graphs which can be manipulated after the fact

```
expected.test.plot$main <- "ETS for Reference and Focal Groups"  
expected.test.plot$legend$top$args$key$text[[1]] <- c('Focal', 'Reference')  
expected.test.plot
```

